

Shortest Path Algorithm with Real World Data

Brad Peterson - Weber State University

Goal: Use Dijkstra's shortest path algorithm with realistic data on a graph data structure that is both fast and avoids memory waste.

This assignment starts with a functioning shortest path algorithm using a 2D array to store graph data. This 2D array is great for a small number of nodes, but wasteful at thousands of nodes as it stores which nodes don't connect with each other. The assignment starts with support for a CSR graph data structure, then adds supporting functions, then modifies Dijkstra algorithm for the CSR graph, and then determines the exact node-by-node path for a shortest path.

CSR Review:

CSR graph storage allows leaving out unconnected and unnecessary node data. For example, the cost for a node from itself to itself is typically zero. But this zero doesn't need to be specified, it's implied. Likewise, a lack of connection between two nodes doesn't need to be specified. For example, observe the waste in the 2D array approach for a graph with 5 nodes and 10 edges

	0	1	2	3	4
0	0	16	99999999	2	3
1	99999999	0	5	99999999	99999999
2	99999999	3	0	99999999	99999999
3	99999999	12	99999999	0	7
4	99999999	10	4	5	0

This graph can be represented in the following CSR data structure:

```
[0][1][2][3][4][5][6][7][8][9][10]
graphWeights: 16 2 3 5 3 12 7 10 4 5
columns: 1 3 4 2 1 1 4 1 2 3
row: 0 3 4 5 7 10
```

These arrays are not intuitive at first glance. To obtain an edge and its weight, such as from node 3 to node 4, perform the following steps. Take the starting node (3), and go to that index in the row array (row[3] = 5). Add one to starting node index value (3+1), and go to that index in the row array, (row[3+1] = 7). The results of 5 and 7 now refer to indexes [5, 7) in the columns array, meaning indexes 5 and 6 but not 7 in the graphWeights and columns array refer to edges coming from node 3. The first edge is column[5] = 1, or in other words, node 3 connects to node 1. That doesn't work, the goal was to find an edge from node 3 to node 4, and this is an edge for node 3 to node 1. So next try column[6] = 4. That 4 the target destination. This edge is for node 3 to node 4. Next, obtain the weight, graphWeights[6] = 7. In summary, an edge exists from node 3 to node 4 with the weight of 7.

Consider going from node 4 to node 3. First, row[4] = 7 and row[4+1] = 10. So, look in [7, 10). The first destination node from node 4 is column[7] = 1, which doesn't match. The second destination node is column[8] = 2, that doesn't work. The third is column[9] = 3. Found it. Get the weight, graphWeights[9] = 5. In summary, an edge exists from 3 to 4 and has a weight of 5.

Consider node 0 to node 2. First, row[0] = 0, and row[0+1] = 3. Look in [0, 3). The destination column[0] = 1, that's not it. The destination column[1] = 3. That's not it. The destination column[2] = 4. That's not it. All

options are exhausted. That means an edge between the two nodes doesn't exist.

Instructions:

Your goal is implementing all TODOs in the assignment. Specifically:

- Initialize arrays in `createCSRArrays()`
 - Allocate five arrays: `graphWeights`, `columns`, `row`, `pathCost`, and `thisCameFrom`.
 - Load the `graphWeights`, `columns`, and `row` arrays with correct values. All graph information is found in the `edges` collection.
 - Initialize every element of the array `pathCost` to `LARGE_NUMBER`. Initialize every element of the array `thisCameFrom` to `-1`.
- Add code for `getEdgeWeight()`
 - Determine the weight for the edge connecting `sourceIndex` and `destIndex` and then return it. If no edge exists, return `LARGE_NUMBER`.
- Add code to `shortestPath()`
 - Replace the `weights[][]` 2D array with a function call `getEdgeWeight()` that uses as arguments the source node and the destination node. The algorithm has three such instances to replace.
 - Every time `pathCost` is updated with a new value for its destination, also update `thisCameFrom`. For example, if the code updates path at node 5 by executing `path[5] = getEdgeWeight(3,5)`, then `thisCameFrom[5] = 3`. Or in other words, the path to node 5 came from node 3.
- Add code to `getPath()`
 - The goal here is following the `thisCameFrom` array from destination back to source. Then return a string indicating the path. The format of the string must be: `"index0 -> index1 -> index2"`. For example, a path of 0 to 4 to 2 to 1 would return the string: `"0 -> 4 -> 2 -> 1"`.
 - The logic for this function is a bit tricky. If the path from node 0 to node 1 requires the path 0 to 4 to 2 to 1, then the logic should start at `thisCameFrom[1]`, see value 2, go to `thisCameFrom[2]`, see value 4, go to `thisCameFrom[4]`, see value 0, and stop as it's the starting index. Note that this logic finds the path from destination to source, so it needs to be reversed prior to building the string.
- Add code to `deleteArrays()`
 - Reclaim the five arrays.

This assignment uses the file `rome99.gr`. The unit tests load all the file data for you into the `edges` collection. CMake should handle placing the file next to the executable correctly regardless of IDE or build process. If not, contact me.